

An Example for Performance Prediction for Map Reduce Applications in Cloud Environments

Iván Carrera, Fabricio Scariot, Cláudio Geyer
Institute of Informatics INF
Federal University of Rio Grande do Sul (UFRGS)
Porto Alegre - RS - Brazil
{ivan.carrera, geyer}@inf.ufrgs.br,
fabricio.scariot@ufrgs.br

Pierre Turin
Ensimag - Grenoble INP
681 rue de la Passerelle - BP 72
F-38402 St Martin d'Hères Cedex, France
Pierre.Turin@ensimag.grenoble-inp.fr

Abstract

One of the advantages of cloud computing is lowering costs to the user by charging only for the computational resources used by the application. In data-intensive applications like Map Reduce, it can be done by using a virtual machine (VM) cluster in the cloud. The goal of this paper is to address the challenge of modelling the behavior of a distributed application running in a Cloud environment and determine the characteristics of the VM cluster that can have the desired performance in the least time. After measuring the time taken by the application and varying parameters of the infrastructure like workload input size and numbers of workers of the cluster, the goal is achieved by finding a model of the execution time which was then applied to predict the execution time for different values of the same variables. Two mathematical models are presented, one for a private cluster and another for a Cloud environment.

1. Introduction

In April 2013, a private company, located in Porto Alegre/Brazil, asked the GPPD/UFRGS laboratory for consultancy to develop and test a Map Reduce application to process large amounts of logs. The logs come from their clients hourly, and have to be processed to feed a Database. The application is meant to be run in a Cloud Computing infrastructure. It is important to know the minimum cluster configuration in terms of quantity of nodes that can meet the demand of this application and the time that will take to accomplish said demand, so the cost of running the application can be minimized. The goal for the consultancy from the GPPD/UFRGS laboratory in this research was to study the behavior of the Map Reduce log processing application as a part of a research that is currently being per-

formed with MapReduce applications over the Cloud. We based our work on a procedure described in [4], which explained a simple but general methodology for modelling the performance of a parallel application running in a Cloud infrastructure, containing 4 steps as follows:

- Defining the parameters that affect the performance of a computer system;
- Perform observations of the application varying the parameters defined in the previous step and taking measures of the performance;
- Propose a mathematical model that relates the performance measures taken with the values of the parameters defined previously, and finally;
- Assess the developed model and test its accuracy.

This methodology will be conducted on several environments to see if it is always accurate. Two environments were chosen: a private cluster and Amazon Web Services' Elastic MapReduce. The AWS platform was chosen by the enterprise, since they are commercial partners.

The remainder of this paper is organized as follows: Section 2 presents the Motivation and background of this work; Section 3 describes how the experiments were conducted and its results; and finally, Section 4 states the conclusions and future work of this research.

2. Motivation

2.1. Map Reduce

Map Reduce is a programming model used to easily implement distributed programs. It was proposed by Google in [5] and there are several implementations. The Map Reduce (MR) model lets the programmer define his Map and Reduce functions, however, it is restrictive enough so that the application can be automatically parallelized, scheduled

and scaled on different machines, which could be physical or virtual. MR applications have a wide range of domains where they can be useful, and that is the reason for research to be done in order to expand it to add more features and correcting some of its features. The Hadoop [2] implementation of Map Reduce was chosen for this case since it can achieve the goal of the application, its ease of use and because it can be easily expanded to add more features, needed by the application.

The Map and Reduce functions are written by the user. For the Map phase, each machine analyses a part of the input, which is divided in pieces called *chunks*. And for the Reduce phase each machine processes a subset of the intermediate results arranged by key.

2.2. Capacity Planning, Performance evaluation and Performance prediction

Capacity Planning is described in [8] as a process where it can be ensured that adequate computer resources will be available to the users of the system to meet future workload demands, and also [9] refers to Capacity Planning, defining it as properly design and size a computer system for a given load condition. In [1], authors present the importance of having Performance models. The developed Performance Model is an analytical model that capture aspects of the system and relate each one by mathematical formulas and/or computational algorithms.

An interesting work in the topic of Map Reduce and Performance Prediction is [10]. In this paper, authors present a cost function that shows a relationship between some characteristics of the Map Reduce application and the time that takes for the application to execute; and also that they present also cost functions for running MapReduce applications on virtual machines in Cloud environments. In [7], authors develop a way for performing Capacity Planning in Cloud computing environments for Map Reduce applications. Authors deal with the problem of determining the virtual machine cluster resources and the Map Reduce configurations to achieve user-defined requirements on execution time and economic cost for a given workload.

2.3. Goal

The goal of this paper is to develop a simple model to predict the performance in terms of execution time of a *log serialization application* running in a private cluster and a cloud computing environment. This model is intended to help the users of the application to perform an accurate capacity planning for the cluster they will use in the cloud computing environment.




Figure 1. Block diagram of the application

3. Experiments

3.0.1. Application The developed Map Reduce application (Fig. 1) receives files containing logs and transforms this input data in a serialized format so the data can then be more easily processed by a Database.

The input data files contain lines of the logs from a CDN server in a text format. Each file contains an hour worth of logs. Later, output files are classified in directories representing the date of the logs. These files contain all the serialized logs for each server, corresponding to the input files.

3.0.2. Performance evaluation To make observations to the application and take measures of performance in terms of execution time, two parameters will be taken into account:

- input files sizes (W - workload) will be tested with values: $U_W = \{1, 5, 10, 20, 25\}$ GB, and
- number of machines that compose the cluster (*workers*) (p)
 - for the private cluster, it will be tested with values: $U_p = \{4, 8, 12, 16\}$;
 - for Amazon, it will be tested with values: $U_p = \{4, 6, 8\}$.

The goal of these measures is to find a function f which can gives us an expression of the execution time t of the application in terms of the number of workers p and the input size workload W : $t = f(W, p)$. Values for parameters W and p were chosen only to have enough data to obtain a linear model of the execution time of the application, since the values of W for the application in a production environment in the cloud can be among those values.

3.1. Experiment description

The set of tested parameters were all the combinations of the elements of U_W and U_p .

3.1.1. Cluster Every combination of parameters (p and W) was tested 10 times. The Map Reduce application was run on a cluster composed of 18 nodes in total. The cluster is located in INF/UFRGS and accessible at `gradep.inf.ufrgs.br` and each node has an Intel Pentium 4 2.79 GHz CPU and 2 GB in RAM.

The Apache recommendations in [3] were used to select an optimal number of Reduce tasks. The number of reduces (N_R) followed this rule: $N_R = p \times n_{CPU} \times 1.75$. Where p is the number of workers, n_{CPU} is the number of CPU cores in the workers, and 1.75 is a value given by Apache.

3.1.2. Cloud services Every combination of parameters (p and W) was tested 2 times. We chose to use AWS' `m1.small` instance type for the master and worker nodes.

3.2. Results

Results are shown in Fig. 2, where we can see that for a fixed number of workers, the execution time increases along with the input workload size. The lines in the graphs represent the given model found as explained in section 3.3.

3.3. Mathematical model

The results were processed with R [6] to complete a linear regression. We used the model:

$$t_{exec} = \beta_0 + \beta_1 W + \beta_2 p + \beta_3 \cdot \frac{1}{p} + \beta_4 W \cdot p + \beta_5 W \cdot \frac{1}{p}$$

Where t_{exec} is the execution time we try to predict, W is the input size in GB, p is the number of workers and the β_i are the coefficients that the linear regression will calculate.

The linear regression for the cluster measures gave the results:

$$\beta_0 = 465.8, \beta_1 = 16.8, \beta_2 = -11.6$$

$$\beta_3 = -714.9, \beta_4 = 2.1, \beta_5 = 569.8$$

$$R^2 = 0.99$$

And for the Amazon measures, it gave:

$$\beta_0 = -2550.3, \beta_1 = -114.1, \beta_2 = 246.1$$

$$\beta_3 = 8616.4, \beta_4 = 6.2, \beta_5 = 1008.1$$

$$R^2 = 0.95$$

The given coefficients of determination R^2 are high, which gives us a strong confidence in it.

Values	Predicted Value	Exp. Value	Error
$W=1, p=6$	401.6	430.7	7.24%
$W=1, p=10$	373.4	374.8	0.38%
$W=1, p=14$	339.7	387.3	14.0%
$W=5, p=6$	899.3	962.1	6.98%
$W=5, p=10$	752.9	734.7	2.42%
$W=5, p=14$	687.9	653.1	5.05%
$W=10, p=6$	1521.4	1442.0	5.22%
$W=10, p=10$	1227.3	1241.0	1.12%
$W=10, p=14$	1123.0	1250.3	11.33%
$W=20, p=6$	2765.6	2526.2	8.66%
$W=20, p=10$	2176.0	2454.4	12.8%
$W=20, p=14$	1993.3	2021.5	1.41%
$W=25, p=6$	3387.7	3392.4	0.14%
$W=25, p=10$	2560.3	2557.0	3.52%
$W=25, p=14$	2428.5	2296.9	5.42%

Table 1. Results for model assessing of the cluster measures

Values	Predicted Value	Exp. Value	Error
$W=1, p=5$	531.0	458.2	13.71%
$W=1, p=7$	488.2	361.5	25.95%
$W=5, p=5$	1061.4	870.2	18.01%
$W=5, p=7$	860.6	780.1	9.26%
$W=10, p=5$	1724.4	1556.9	9.71%
$W=10, p=7$	1326.2	1186.7	10.51%
$W=20, p=5$	3050.3	2826.9	7.33%
$W=20, p=7$	2257.2	2928.5	29.74%
$W=25, p=5$	3713.3	3556.4	4.23%
$W=25, p=7$	2722.7	2414.6	11.32%

Table 2. Results for model assessing of the Amazon measures

3.4. Assessing the Model

After the mathematical model was proposed, it had to be assessed with experimental results. The same values of W were used, being $W = \{1, 5, 10, 20, 25\}$ GB. The new values to assess the model were for p , being $p = \{6, 10, 14\}$ nodes for the cluster experiments and $p = \{5, 7\}$ nodes for the Amazon experiment. Experiments in this part were conducted as explained in section 3.1, and the average values are described in Tables 1 and 2.

From the cluster results, shown in Table 1, we can see that the model was good enough to predict the execution time for the conditions expressed in the Values column, with

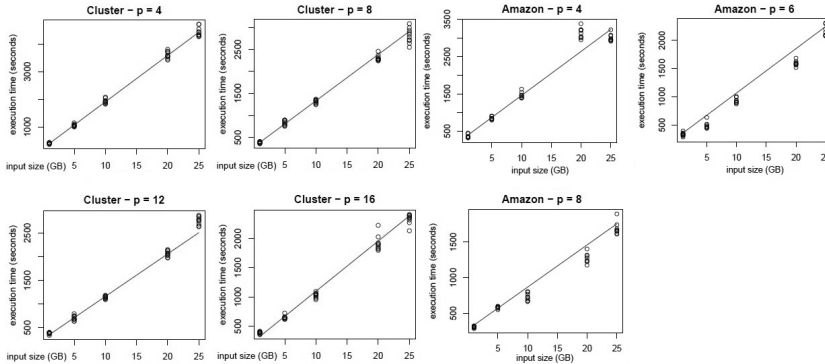


Figure 2. Results of execution time (in seconds) against input size of the workload (in GB) of the experiments, with indicated number of workers running the Map Reduce application

an error no bigger than 14% of the predicted value. From the Amazon results, shown in Table 2, we can see that most errors are low except the two of them that range near from 35%.

With these results we say that the models are accurate and valid for execution time predictions with this application.

4. Conclusions and Future Work

Based on the results of section 3.4, we see that a mathematical model can be obtained and applied to predict the execution time of a Map Reduce application on different executing environments.

Using the methodology described in [4] helped to model the behaviour of a parallel application in a private cluster environment. So, a future work will be to research about the variance of the results, in a field known as *Performance Debugging*.

5. ACKNOWLEDGEMENTS

This work was made with the support of the *Programa Estudantes-Convênio de Pós-Graduação PEC-PG, of CAPES/CNPq - Brazil*.

References

- [1] V. A. Almeida and D. A. Menascé. Capacity planning an essential tool for managing web services. *IT professional*,

4(4):33–38, 2002.

- [2] Apache, 2013. Welcome to Apache Hadoop! <http://hadoop.apache.org/> accessed on 24/07/2013.
- [3] Apache, 2013. HowManyMapsAndReduces - Hadoop Wiki <http://wiki.apache.org/hadoop/HowManyMapsAndReduces> accessed on 24/07/2013.
- [4] I. Carrera and C. Geyer. Impressionism in cloud computing - a position paper on capacity planning in cloud computing environments. In *Proceedings of the 15th International Conference on Enterprise Information Systems ICEIS - Vol. 2*, pages 333 – 338. INSTICC, 2013.
- [5] J. Dean and S. Ghemawat. Mapreduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.
- [6] R. Foundation, 2013. The R Project for Statistical Computing <http://www.r-project.org/> accessed on 24/07/2013.
- [7] H. Herodotou, F. Dong, and S. Babu. No one (cluster) size fits all: automatic cluster sizing for data-intensive analytics. In *Proceedings of the 2nd ACM Symposium on Cloud Computing*, page 18. ACM, 2011.
- [8] R. Jain. *The art of computer systems performance analysis*, volume 182. John Wiley & Sons Chichester, 1991.
- [9] D. A. Menascé, V. A. Almeida, L. W. Dowdy, and L. Dowdy. *Performance by design: computer capacity planning by example*. Prentice Hall Professional, 2004.
- [10] F. Tian and K. Chen. Towards optimal resource provisioning for running mapreduce programs in public clouds. In *Cloud Computing (CLOUD), 2011 IEEE International Conference on*, pages 155–162. IEEE, 2011.